

# Rylee: Creating a Deployable Human-Like Chess Engine to Enhance the Learning Experience

---

Ethan Gee & Nate Stott

# Agenda

---

1. Introduction & Motivation
2. Methodology
3. Experiments
4. Conclusions & Future Work

# Introduction - Problem Definition

---

## Maia Problems

- Traditional chess engines maximize the chances of winning
- Predicting a move a human would play does not mean finding the best move
- **Goal:** Replicate human play

## Rylee Problems

- Maia requires **large GPUs** to run and train
- Maia can only run on high end machines
- Maia was not trained to play chess openings
- **Goal:** Make a **practical & deployable** version of Maia

# Introduction - Motivation

---

## Why human aligned AI matters

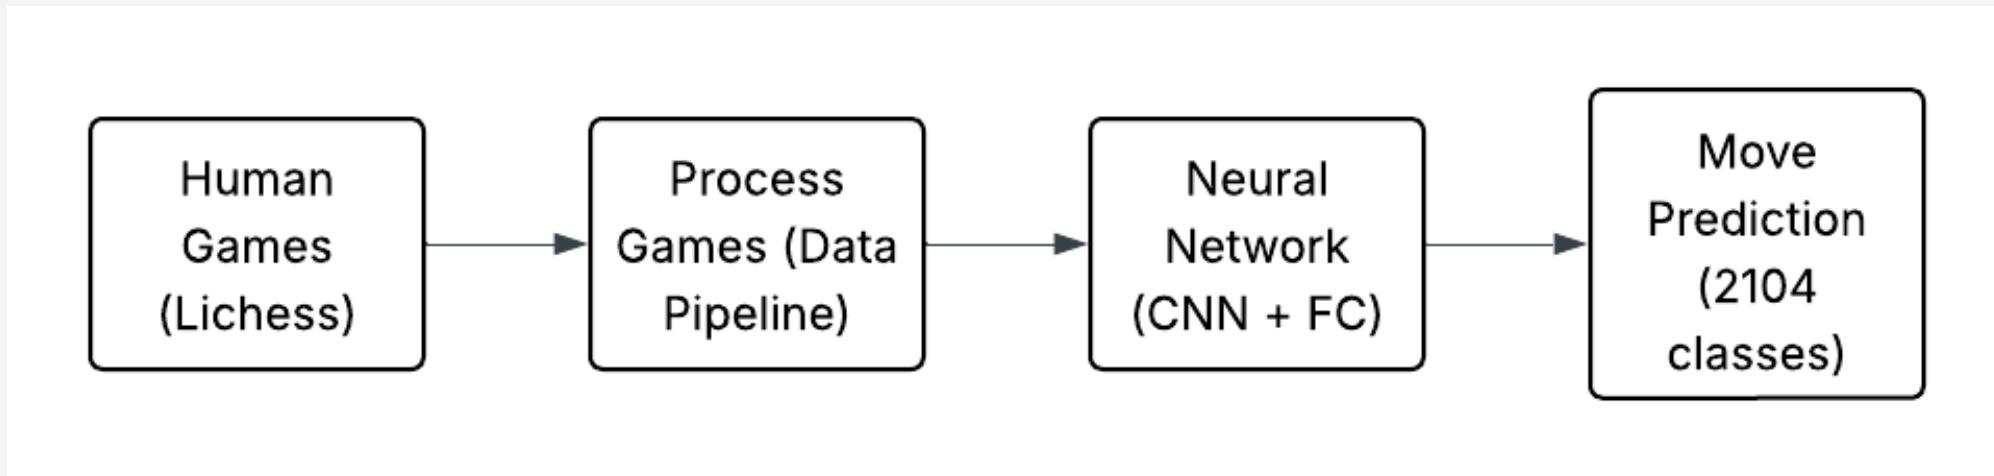
- Traditional engines play chess differently making it difficult for humans to learn from
- Attenuating does not **mimic human play**
- Human aligned engines creates more realistic **training partners**
- Example: Chess students can practice with Rylee on their school chromebook to advance their chess skills
- Broader applications: Collaborative decision-making, Education, etc

## How Rylee extends Maia

- **Edge deployment** Raspberry Pi, Chromebooks
- **Includes openings** first 10 moves
- **No game filtering** include all game types (classical, blitz, etc)
- **Unified model with a higher range** 700-2500 ELO

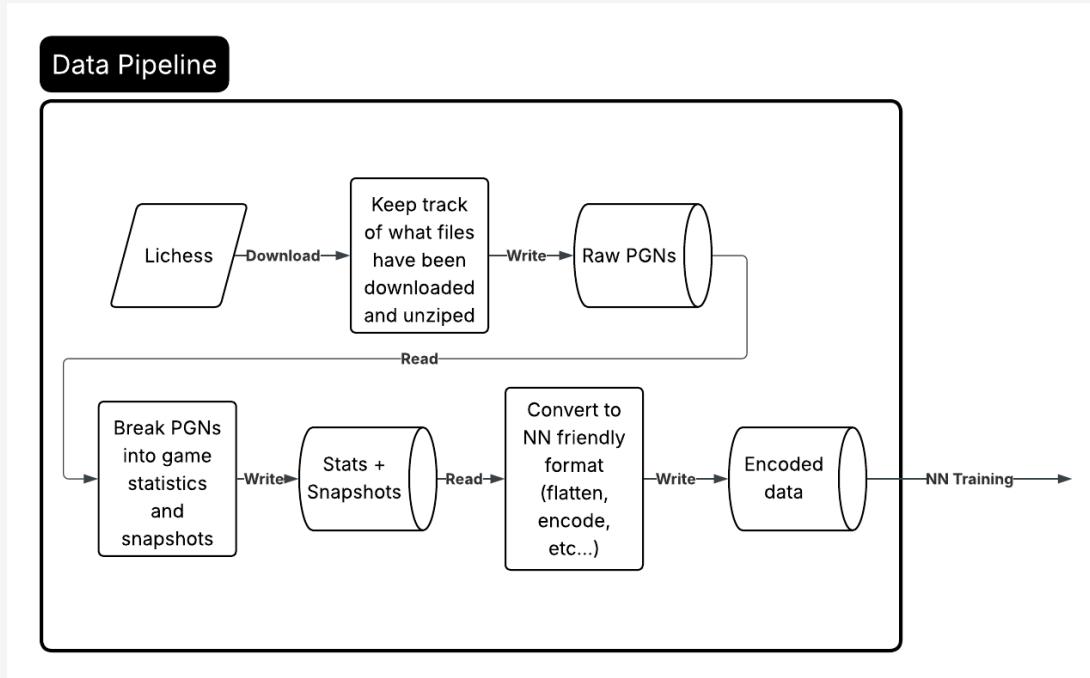
# Methodology - Proposed Solution

*We hypothesize we can maintain similar performance, and add features to the Maia model while significantly reducing model size.*



1. Pull data from Lichess
2. Preprocess games
3. Feed data into NN
4. Predict moves

# Methodology - Data Pipeline



- **Download** .zst files from Lichess
- **Extract** PGNs
- **Split** into individual games
- **Convert** to board snapshots
- **Extract** ELO and result metadata
- **Encode** board as 8x8x12 tensors

# Methodology - Theories

---

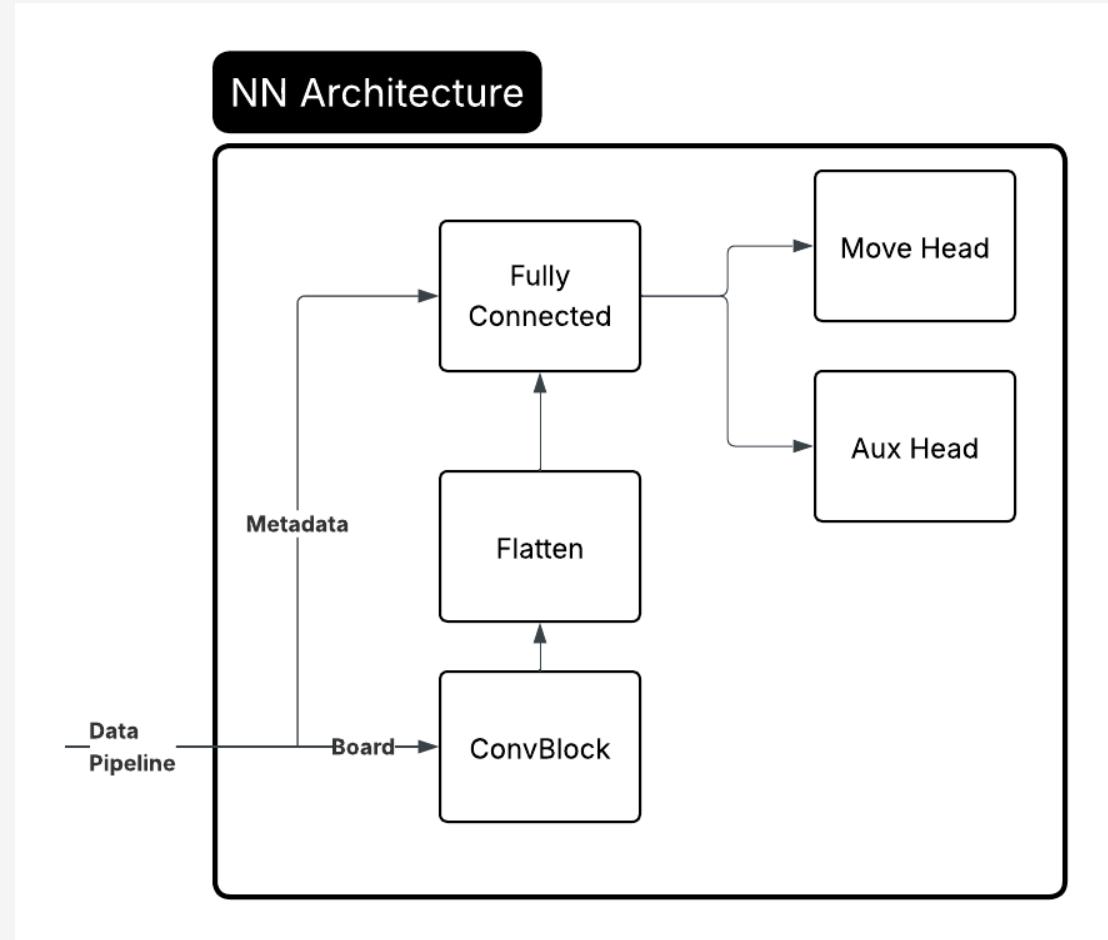
## Why CNNs work well for chess

- Chess boards are **spatially related** (knight is better if its in the middle)
- Humans evaluate through pattern recognition
- CNNs excel at **spatial pattern recognition**

## Model size vs performance

- Increasing model size has exponentially diminishing returns
- We are hoping we can decrease the size of the Maia model while still keeping high accuracy

# Methodology - Neural Network



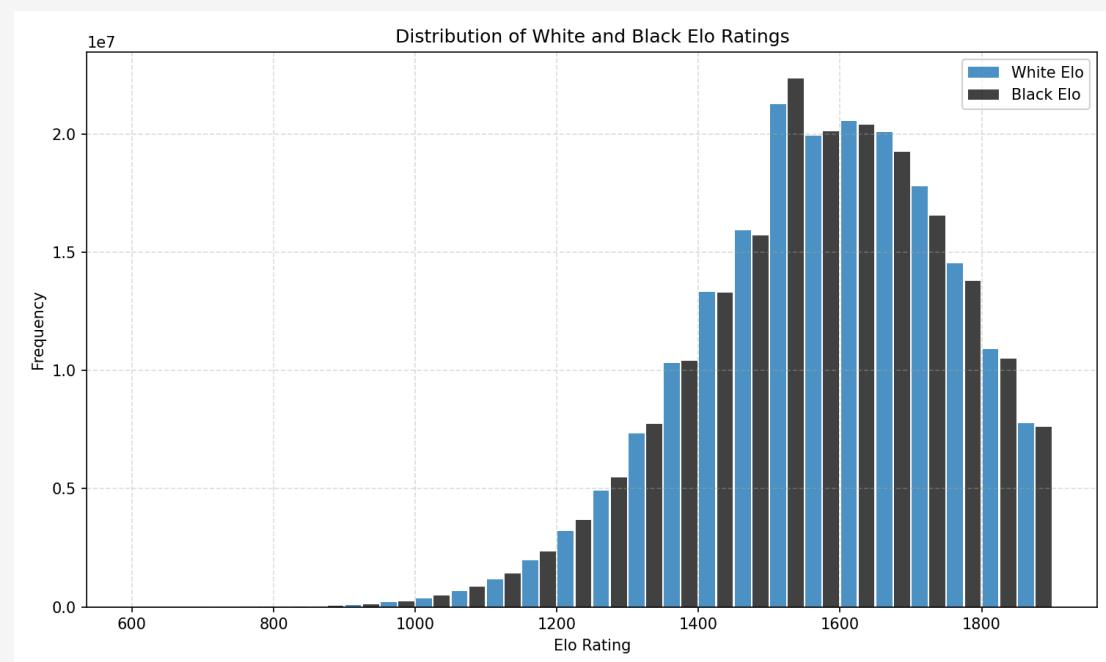
- **Input:** Board(8x8x12) + Metadata(4)

- **Conv Layers:** 6 64x8x8 filters, ReLU
- **Fully Connected:** 4100  $\rightarrow$  512  $\rightarrow$  32
- **Output Heads:** Move (2104) + Auxiliary (2104)
  - Move Head = predicted chess move
  - Aux Head = predicted legal chess moves
  - 2104 is the number of legal moves
- **Loss:** CrossEntropy (moves) + BCE (valid moves)
- **Optimizer:** Adam
- **Hyperparameter Search:** Random search

# Experiments - Dataset

- **Source:** Lichess Open Database
- **Games:** 15,167 human-rated games
- **Snapshots:** 1 million board states
  - including openings
  - including all game types
- **Action Space:** 2,104 legal move classes
- **Time Span:** January 2013
- **Expanded Capabilities:** Covering 5x the amount of data

| Split      | Percentage | Snapshots |
|------------|------------|-----------|
| Training   | 80%        | 800,000   |
| Validation | 10%        | 100,000   |
| Test       | 10%        | 100,000   |



# Experiments - Baselines

---

| Baseline Model | Description   |
|----------------|---|
| Random         | Random legal move selection                           |
| Random Forest  | Nothing that simple should work that well - Ethan Gee |
| Stockfish 15   | Traditional chess engine                              |
| Leela 4200     | Neural chess engine                                   |
| Maia-1 1500    | Human aligned prediction model                        |

# Experiments - Architecture

---

## Small Fully Connected Model

- A Small model that had a similar architecture to StockFish
- 8 fully connected layers of 32 neurons

## Convolutional Model

- Combination of Convolution and fully connected to mirror human cognition

## Convolution with Auxillary Head

- Added an auxillary head that determines legal moves to instill better game understanding

## Experiments - Evaluation Metrics

---

- **Top-1 Accuracy**: Predicted move matches actual human move
- **Top-5 Accuracy**: Actual move in top 5 predictions. This is good for a more generalized alignment.

# Experiments - Comparisons

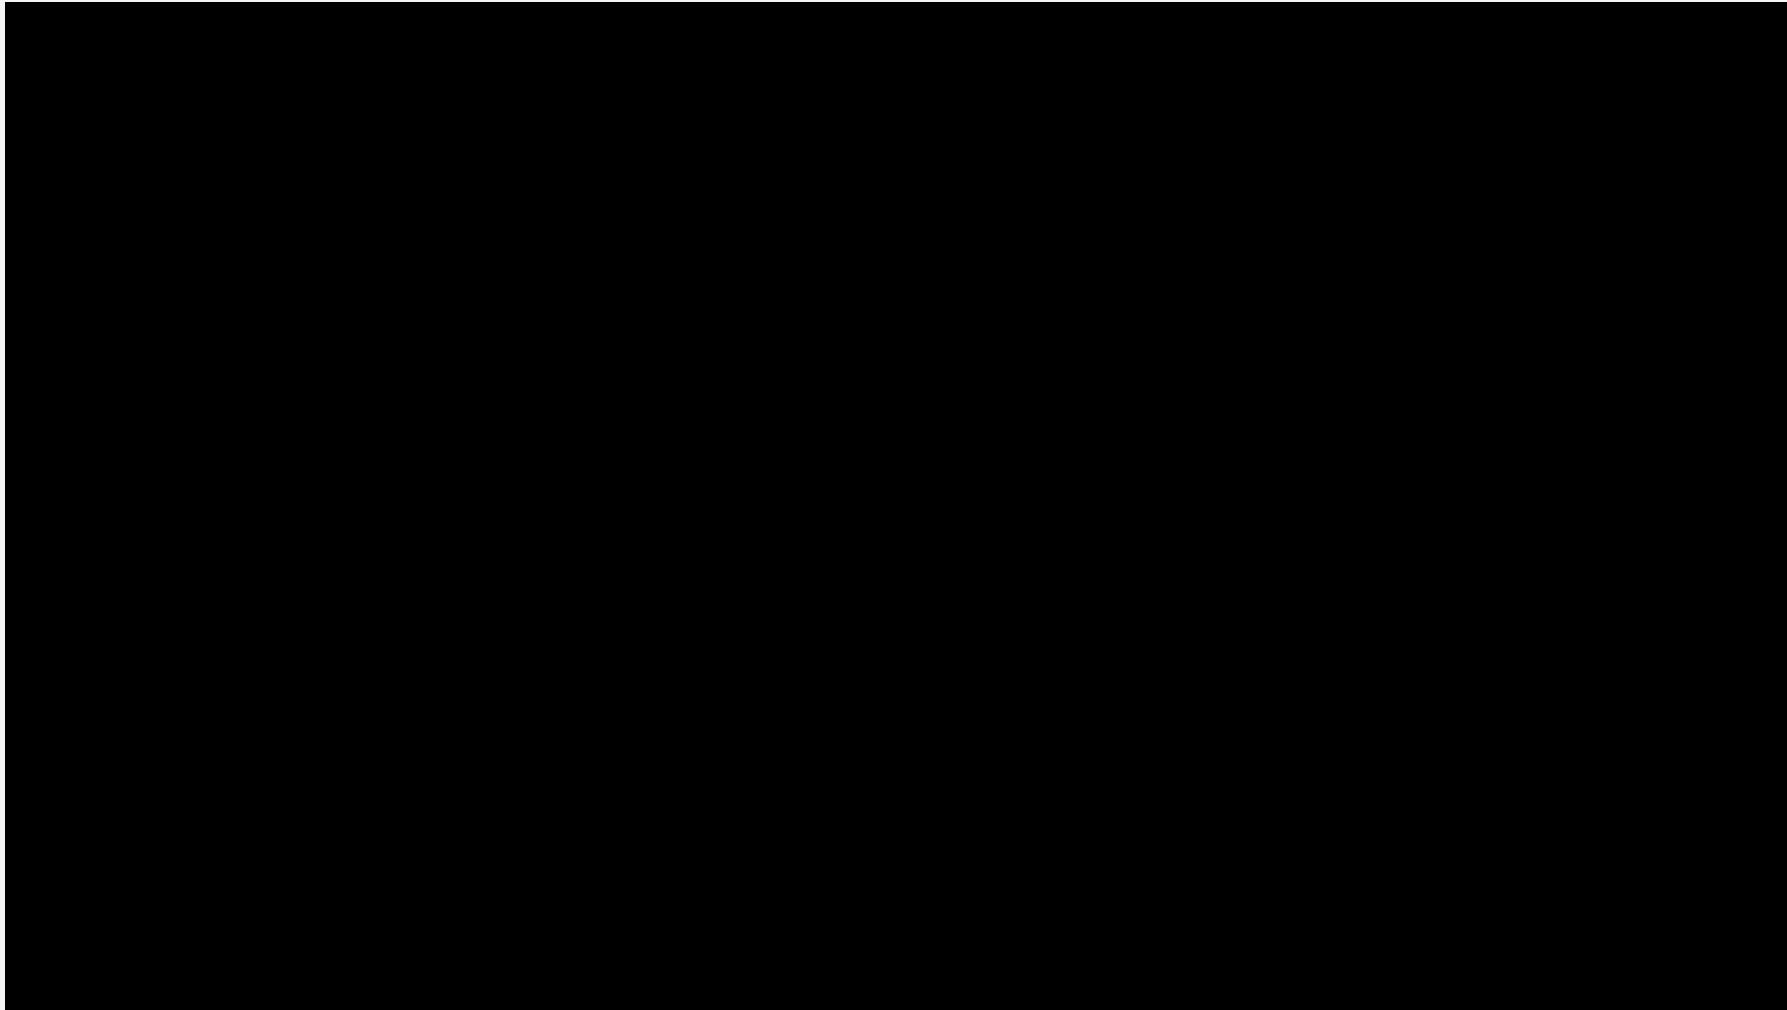
---

| Method                       | Top-1 Accuracy |
|------------------------------|----------------|
| Random                       | 6%             |
| Random Forest                | 13%            |
| Stockfish 15                 | 40%            |
| Leela 4200                   | 44%            |
| Maia1 1500                   | 51%            |
| Rylee FC                     | 3.5%           |
| Rylee Conv                   | 23.5%          |
| Rylee Conv with Aux          | 25%            |
| Rylee Conv with Aux Filtered | 35%            |

- Rylee has 800,000 parameters vs Maia's 25 Million
- No filtering by game type (classical, blitz, etc) to capture broader human play patterns
- We include games with mixed skill levels to better reflect general human behavior
- 15,000 games vs. Maia's 169 million games
- Maia was Trained on two A100 80Gb GPUs vs Rylee being trained on a Edge Device

## Experiments - Deployment

---



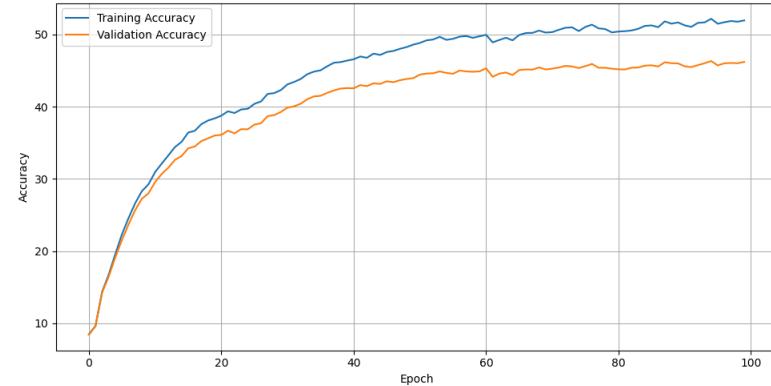

We have a tkinter gui. Here is it running on a crappy laptop. Black is Rylee and White is Stockfish.

# Conclusions - Discussions

| Metric                  | Training | Validation |
|-------------------------|----------|------------|
| Loss                    | 0.0152   | 0.0164     |
| Top-1 Accuracy          | 27%      | 25%        |
| Top-5 Accuracy          | 53%      | 51%        |
| Top-1 Accuracy Filtered | 36%      | 35%        |
| Top-5 Accuracy Filtered | 54%      | 53%        |

- Strong generalization between training and validation metrics. Model captures key human decision-making patterns.
- Rylee required around 1.5 hours of preprocessing and 2-3 days of training

- Maia required 8 days of preprocessing and 3-4 weeks of training



# Conclusions - Future Work

---

## Model Improvements

- Add data augmentation (board flips and rotations) to improve robustness
- Time parameter to better address time based decision making
- Cross Validation
  - Maia was not able to do this because of the size of the dataset

## Additional Features

- **ELO Prediction:** Estimate player rating from move patterns to quickly adapt to player skill
- **Human vs Bot Discriminator:** Detect engine-like play
- **Blunder Detection:** Identify major mistakes for analysis

## Conclusions - Summary

---

- Rylee mimics human chess behavior using a model that is 30x smaller than Maia
- Achieves competitive accuracy (25-35% Top-1, 51-54% Top-5) despite using significantly less data, compute, and expanding capabilities, and data variety
- Efficient data pipeline and compact architecture make Rylee deployable edge devices such as Chromebooks and Raspberry Pis
- Generalizes well across training and validation datasets, indicating a healthy fitting of human chess playing patterns
- Demonstrates the feasibility of edge-deployable, human-aligned AI for education applications

# References

---

## Primary Works

- McIlroy-Young et al. (2020). "Aligning Superhuman AI with Human Behavior: Chess as a Model System." KDD 2020.
- Tang et al. (2024). "Maia-2: A Unified Model for Human-AI Alignment in Chess." NeurIPS 2024.
- McIlroy-Young et al. (2021). "Detecting Individual Decision-Making Style: Exploring Behavioral Stylometry in Chess." NeurIPS 2021.

## Rylee Repo:

<https://github.com/EthanDGee/ryleeeeeeeeeeee>

## Data & Tools

- Lichess Open Database: <https://database.lichess.org/>
- Stockfish Chess Engine: <https://stockfishchess.org/>
- Leela Chess Zero: <https://lczero.org/>
- Maia Chess Project: <https://maiachess.com/>
- PyTorch (Paszke et al., 2019): <https://pytorch.org/>
- python-chess library (Moskopp, 2014): <https://github.com/niklasf/python-chess>

# Terminology

---

- **ELO Rating** - Numeric chess player skill score used to represent player strength (500 beginner, 1500 intermediate, 2500 expert)
- **Board Snapshot** - A single chess board state (imagine taking a picture of the board every time a player makes a move, each of those pictures should be a board snapshot)
- **Action Space (2104 moves)** - Fixed index set representing all possible legal chess moves
- **Auxiliary Head** - Secondary output predicting legal moves to guide the main move head. Meant to strengthen legal move connections/predictions.
- **Opening Phase** - The first 10ish moves of the game
- **Blunder** - An objectively terrible chess move
- **Lichess Dataset** - Large open database of real human chess games
- **Human-Aligned Model** - Predicts human-like moves rather than optimal engine moves

# Questions?

---

Rylee: Creating a Deployable Human-Like Chess Engine to Enhance the Learning Experience

Ethan Gee & Nate Stott